

Comparing Frontier Large Language Models for Cybersecurity Operations:

A MITRE ATT&CK and ATLAS Analysis of Anthropic Claude, OpenAI GPT-5, and Google Gemini for SecOps

Johnbosco Chizitem Ibeneme

Cybersecurity Specialist Intern, disruptiveOps

May 2026

Abstract

Large Language Models (LLMs) are being used by both the offensive side and the defensive side in the field of cybersecurity. Vendors are now productizing their own security-focused versions of LLMs. This paper will compare three of these models — Anthropic Claude (Claude Security & Claude Mythos — a restricted version of Claude); OpenAI GPT-5 (Codex Security & Trusted Access for Cyber versions); and Google Gemini for SecOps (built on the SecLM platform), using five evaluation metrics and through comparison to the MITRE ATT&CK framework and the MITRE ATLAS framework. Five ATT&CK techniques and four ATLAS techniques were mapped to each of the models. The results of this research show that there is a gap in capabilities versus policies, where the ability to deploy the LLM can be controlled more by the actual deployment environment/controls rather than the benchmark scores of the LLM. There is also an asymmetrical approach to how the vendors specialize in different areas of operation. Finally, all of the models deployed for defensive purposes create a target for attacks based on previously published methods. These findings suggest four recommendations for security organizations when considering using frontier LLMs in a production capacity.

Keywords: large language models, red teaming, blue teaming, MITRE ATT&CK, MITRE ATLAS, prompt injection, AI security, cybersecurity.

1. Introduction

Cybersecurity is at an inflection point: the gap between vulnerability discovery and exploitation is narrowing as AI strengthens attacker capabilities. Anthropic noted in April 2026, while announcing Claude Security, “AI cybersecurity capabilities are advancing fast,” framing the moment as a milestone, one that gives defenders access to cutting-edge capabilities. The evidence supports the claim: OpenAI’s Aardvark agent achieved a 92% detection rate on benchmark vulnerability repositories, and Anthropic’s restricted Claude Mythos identified 271 vulnerabilities in Mozilla Firefox. As vendors continue to produce large language models (LLMs) for defensive and offensive cyber operations, the cyber industry is faced with a gap; vendors only address what is possible, but do not discuss product fit for the operational workflow. This paper addresses that gap.

1.1. Dual-Use Dilemma

This capability inflation is a double-edged sword. Abuadbbba and colleagues, in their 2025 paper on LLMs in cybersecurity, argue that LLMs simultaneously strengthen offensive and defensive operations. Offensively, LLMs can plan attacks, craft phishing emails, simulate adversarial behaviors, and generate exploit code. Defensively, they can synthesize threat intelligence, support root cause analysis, and streamline security documentation. The dual-use attribute of these LLMs means that for every defensive capability there is an equal offensive capability, and by deciding who gets access to these models, what requests are blocked, and how safety checks are carried out, vendors indirectly balance offensive and defensive capabilities.

1.2. Contribution

The dual-use dilemma has received a lot of attention, but comparative analysis has received little. Vendors have failed to address product fit for different operational workflows or address which models best withstand attacks directed at LLMs themselves. This paper compares three cutting-edge LLMs — Anthropic Claude, OpenAI GPT-5, and Google Gemini for SecOps — across the MITRE ATT&CK red-team and ATLAS adversarial frameworks. The analysis produces three findings: a capability-versus-policy gap in which deployment restrictions determine what users are actually allowed to do more than benchmark scores; differing areas of specialization in which the three vendors focus on different practical capabilities; and a defender-as-target paradox, where every defensively utilized LLM is itself an attack surface.

1.3. Roadmap

This paper proceeds as follows. Section 2 evaluates the frameworks used for analysis in this paper, including NIST AI RMF, MITRE ATT&CK, MITRE ATLAS, and the OWASP Top 10 for LLMs, and examines the 2025–2026 threat landscape from both academic and practitioner sources. Section 3 classifies the three

LLMs, spanning model versions, security-specialized variants, use cases, and access tiers. Section 4 presents the methods of evaluation, including the five-criterion comparison framework summarized in Table 1. Section 5 builds the central reasoning behind the paper: a mapping of five ATT&CK techniques and four ATLAS techniques to each of the three LLMs, in Tables 2 and 3. Section 6 enumerates the findings and their implications for security organizations. Section 7 concludes with limitations and a forward trajectory.

2. Background and Related Work

2.1. Red, Blue, Purple Teaming and the Dual-Use Frame

In the context of practitioner literature, red teaming is a security assessment method in which an adversary-emulating team simulates real-world attacks to identify weaknesses across an organization’s people, processes, and technology. On the defensive side, blue teams monitor, detect, and respond to such activity. Purple teams facilitate collaboration between the two. Red-team engagements typically follow a sequence: planning, reconnaissance, initial access, persistence, privilege escalation, lateral movement, objective completion, and reporting. LLMs are now in a unique position to address both edges of the proverbial sword, a duality that Abuadbba and colleagues describe as innovative and risky.

2.2. Frameworks

This paper’s analysis is supported by four frameworks. The NIST AI Risk Management Framework (AI RMF 1.0, NIST AI 100-1), published in January 2023 and supported by the Generative AI Profile (NIST AI 600-1), provides guidance on managing risks in AI and generative AI systems. The Critical Infrastructure Profile concept note, published in April 2026, enhances the framework using sector-focused risk management practices. MITRE ATT&CK provides a standardized classification of adversarial methods and techniques observed in the real world. MITRE ATLAS classifies adversarial tactics used against AI and ML systems; AML.T0051 (LLM Prompt Injection) is the most relevant attack technique in this analysis. The OWASP Top 10 for LLM Applications (2025) names Prompt Injection (LLM01), Sensitive Information Disclosure (LLM02), and Excessive Agency (LLM06) as the highest-priority risks.

2.3. Recent Threat Landscape

A rapid expansion in the threat landscape of production LLMs has been observed. A research team led by Aonan Guan at Johns Hopkins University, in April 2026, reported a “comment-and-control” prompt-injection method that exploited three production LLM agents in GitHub Automation Workflows: Anthropic Claude Code Security Review, Google Gemini CLI Action, and Microsoft GitHub Copilot. This attack disclosed Anthropic’s API key, GitHub access tokens, and the Gemini API key. Anthropic, Google, and

Microsoft all paid bug bounties; no CVEs or public advisories were issued. Months earlier, researchers at Varonis reported the Reprompt attack against Microsoft Copilot, a single-click data exfiltration chain that exploits indirect prompt injection. Other related vulnerabilities — GeminiJack in Gemini Enterprise, CellShock in Claude for Excel, and Lies-in-the-Loop affecting Claude Code and Microsoft Copilot Chat — surfaced during the same period. The pattern is clear: every LLM analyzed in this study is also a documented attack surface.

3. The Three LLMs

3.1. Anthropic Claude

As of May 2026, Anthropic’s flagship model is Claude Opus 4.7, with a security-focused variant, Claude Security, released on April 30, 2026, in public beta for Claude Enterprise customers. Claude Security scans repositories, directories, or branches for weaknesses and generates explanations, severity confidence scores, replication steps, and patches without requiring API integration or agent deployments. Claude Mythos, a more restricted model, is only accessible through Anthropic’s closed Project Glasswing; in security testing it reportedly identified 271 vulnerabilities in Mozilla Firefox. Technology partners CrowdStrike, Palo Alto Networks, SentinelOne, Trend.ai, and Wiz are integrating Opus 4.7 into their security platforms; service partners Accenture, BCG, Deloitte, Infosys, and PwC aid enterprise deployments.

3.2. OpenAI GPT-5

OpenAI’s flagship model GPT-5 was released in August 2025, with a successor model GPT-5.5 released on April 23, 2026, and classified at “high” cyber capability according to OpenAI. The security variant, originally named Aardvark and renamed Codex Security in March 2026, is powered by GPT-5 under the hood and is deployed as an agentic researcher that monitors source-code commits, builds a project threat model, validates exploitability in sandboxed environments, and generates patches through OpenAI Codex. During benchmark testing on “golden” repositories with known vulnerabilities, Aardvark delivered a 92 percent detection rate. Aardvark has additionally produced over ten publicly assigned CVE identifiers for issues discovered in open-source software. GPT-5.4-Cyber and GPT-5.5-Cyber are available only to vetted organizations through OpenAI’s Trusted Access for Cyber (TAC) program; these models offer lower refusal boundaries for cybersecurity-focused offensive work and add binary reverse engineering capabilities.

3.3. Google Gemini for SecOps

Google, rather than deploying a general-purpose model like Anthropic and OpenAI, has instead deployed Gemini in Google SecOps on the SecLM platform, a specialized base model trained on security blogs,

threat-intelligence reports, YARA and YARA-L detection rules, SOAR playbooks, and malware analysis. It operates inside Google Security Operations, an enterprise SIEM that supports text-based queries. A Triage and Investigation Agent handles routine alert reviews; community-documented use cases include threat-actor research workflows that cross-reference threat intelligence with documented indicators of compromise. Gemini for SecOps is included in the Google SecOps Enterprise package.

4. Methodology

4.1. Evaluation Criteria

Five evaluation criteria adapted from standard security-product evaluation practice are applied in this analysis. Capability fit checks whether the LLM returns output that is useful for specific cybersecurity workflows. Accuracy and hallucination risk evaluate the rate of confabulation and the resilience of the design against that failure mode. Data handling and confidentiality assess the privacy controls and certifications that determine whether the LLM can properly handle sensitive data. Operational integration measures how readily the LLM plugs into existing security stacks: SIEMs, EDR tools, and CI/CD pipelines. Cost and accessibility evaluate the affordability for a new organization, including pricing tiers and vetting requirements. Table 1 summarizes the three LLMs across these five criteria using a five-level scale (HIGH, MED-HIGH, MED, MED-LOW, LOW).

Criterion	Anthropic Claude	OpenAI GPT-5	Google Gemini for SecOps
Capability Fit	HIGH (code-centric); MED-LOW (SOC integration)	HIGH (92% benchmark; agentic pipeline)	HIGH (SOC workflows); LOW (vulnerability discovery)
Accuracy / Hallucination Risk	MED-HIGH (some documented prompt-injection gaps)	HIGH (sandboxed exploit validation; published benchmark)	MED-HIGH (SecLM domain specialization; no public benchmark)
Data Handling / Confidentiality	HIGH (Cyber Verification Program; transparent disclosure)	HIGH (TAC vetting; Preparedness Framework)	HIGH (CMEK; VPC Service Controls; Vertex AI processing)
Operational Integration	MED-HIGH (GitHub-native; partner ecosystem)	MED-HIGH (Codex Security GitHub-native)	HIGH (native SIEM/SOAR; 300+ tool integrations)
Cost / Accessibility	MED (Enterprise tier; Mythos closed)	MED-HIGH (broad entry options; OSS program; \$10M grant)	MED-LOW (Enterprise tier required; no free tier)

Table 1. Comparative Evaluation Matrix of Three Frontier LLMs Across Five Criteria.

4.2. Mapping Methodology

For the mapping in Section 5, five MITRE ATT&CK techniques were selected. They span the offensive kill chain from reconnaissance through lateral movement: T1566 (Phishing), T1595 (Active Scanning), T1059 (Command and Scripting Interpreter), T1190 (Exploit Public-Facing Application), and T1078 (Valid Accounts). Four MITRE ATLAS techniques for AI and ML systems were also used: AML.T0051.000 (Direct Prompt Injection), AML.T0051.001 (Indirect Prompt Injection), AML.T0024 (Exfiltration via ML Inference API), and AML.T0086 (Exfiltration via AI Agent Tool Invocation). Integration depth, capability, and access tier were used to determine the best fit, rather than benchmarks alone.

5. Analysis

5.1. ATT&CK Mapping Overview

Table 2 maps five MITRE ATT&CK techniques to the three LLMs across red-team and blue-team uses. For each technique, the best-fit LLMs are identified separately for offensive and defensive applicability, along with the operational rationale supporting each selection. The same LLM rarely fits both operations, a consequence of the differing operational specialization of the models.

ID	Technique	Red-Team Best Fit	Blue-Team Best Fit
T1566	Phishing	GPT-5.4-Cyber (Trusted Access for Cyber)	Gemini in Google SecOps (native email-log parsers)
T1595	Active Scanning	Claude Opus 4.7 (long-context scanner output reasoning)	Gemini in Google SecOps (UDM-native ingestion)
T1059	Command and Scripting Interpreter	GPT-5.4-Cyber / GPT-5.5-Cyber (binary reverse engineering)	Gemini in Google SecOps (SecLM trained on YARA-L)
T1190	Exploit Public-Facing Application	Codex Security (92% benchmark); Claude Mythos (271 Firefox vulns)	Claude Security; Codex Security
T1078	Valid Accounts	Claude Opus 4.7 (whole-codebase IAM reasoning)	Gemini in Google SecOps (Triage and Investigation Agent)

Table 2. MITRE ATT&CK Techniques Mapped to Best-Fit LLMs for Red-Team and Blue-Team Operations.

5.2. Workflow-by-Workflow Analysis

For T1566 (Phishing), GPT-5.4-Cyber emerges as the best-fit red-team option due to its explicit lowering of refusal boundaries for legitimate offensive work under the Trusted Access for Cyber program; Gemini in Google SecOps is best-fit on the blue side because its native parsers for Mimecast, Proofpoint, and Microsoft 365 enable analyst-driven email-log triage. For T1595 (Active Scanning), Claude Opus 4.7's

long-context reasoning over scanner output makes it the best red-team fit, while Gemini’s UDM-native data pipeline handles defensive ingestion. For T1059 (Command and Scripting Interpreter), GPT-5.4-Cyber and GPT-5.5-Cyber are best-fit offensively, with binary reverse engineering now in scope under TAC; SecLM’s training on YARA-L detection rules makes Gemini best-fit defensively. For T1190 (Exploit Public-Facing Application), Codex Security and Claude Mythos lead offensively (with 92 percent benchmark detection and 271 Firefox vulnerabilities, respectively), while Claude Security and Codex Security lead defensively — the same model architectures repackaged for opposite operational contexts. For T1078 (Valid Accounts), Claude’s whole-codebase reasoning fits IAM analysis offensively; Gemini’s Triage and Investigation Agent handles login telemetry defensively.

5.3. The Pattern in Table 2

A clear pattern emerges from Table 2. Gemini in Google SecOps wins every blue-team row, as operational integration with the SIEM and SOAR stack forms the entire product. Claude wins in rows requiring deep code reasoning over the target, including scanner output, IAM configurations, and scripting logic. GPT-5 wins the red-team rows that require access-controlled offensive capability, since its Trusted Access for Cyber program is the only one among the three that explicitly lowers refusal boundaries for legitimate offensive work.

5.4. ATLAS Mapping Overview

Table 3 shifts the analytical perspective. Rather than viewing the three LLMs as tools that assist with adversary techniques, the table documents attacks targeting the LLMs themselves. Four MITRE ATLAS techniques are mapped to specific incidents observed in 2025–2026. Each row presents an incident affecting at least one of the three LLMs in the comparison, alongside the corresponding vendor response.

ID	Technique	Documented Incident (2025–2026)	Vendor Response
AML.T0051.000	Direct Prompt Injection	Jailbreak techniques documented against all three vendors; stakes highest for cyber-permissive variants (GPT-5.4/5.5-Cyber, Claude Mythos)	Cyber Verification Program (Anthropic); TAC vetting (OpenAI); SecLM domain restriction (Google)
AML.T0051.001	Indirect Prompt Injection	Comment-and-control attack on Claude Code Security Review, Gemini CLI Action, GitHub Copilot (Guan et al., April 2026); GeminiJack; CellShock	Bug bounties paid (\$100, \$1,337, \$500); no CVEs or public advisories; Anthropic updated documentation
AML.T0024	Exfiltration via ML Inference API	Reprompt attack on Microsoft Copilot (Varonis, January 2026):	Microsoft patched; underlying technique class unresolved

ID	Technique	Documented Incident (2025–2026)	Vendor Response
		URL-parameter injection chained with guardrail bypass	
AML.T0086	Exfiltration via AI Agent Tool Invocation	Comment-and-control attack leaked GitHub access tokens, Anthropic API key, Gemini API key via tool invocations	Bug bounties paid; Lies-in-the-Loop exploit shows HITL safeguards can be bypassed

Table 3. MITRE ATLAS Techniques Mapped to Documented Attacks Against the Three LLMs.

5.5. Attack-by-Attack Analysis

For AML.T0051.000 (Direct Prompt Injection), academic and security research literature has documented jailbreak techniques targeting all three vendors; the stakes are highest for OpenAI’s cyber-permissive variants and Anthropic’s Claude Mythos, as the lowered refusal boundary amplifies the consequences of a successful jailbreak. For AML.T0051.001 (Indirect Prompt Injection), the April 2026 comment-and-control attack reported by Guan and colleagues stands as the most prominent contemporary incident, simultaneously compromising Claude Code Security Review, Gemini CLI Action, and GitHub Copilot through malicious instructions embedded in pull-request titles and HTML comments. For AML.T0024 (Exfiltration via ML Inference API), the Reprompt attack on Microsoft Copilot, disclosed by Varonis in January 2026, chains URL-parameter injection with guardrail bypass to allow single-click data exfiltration. For AML.T0086 (Exfiltration via AI Agent Tool Invocation), the comment-and-control attack again offers direct empirical evidence; the agents leaked GitHub access tokens, the Anthropic API key, and the Gemini API key through their connected tool invocations.

5.6. The Defender-as-Target Paradox

The overall trend of data in Tables 2 and 3 supports the main point that this paper makes. While the same LLMs are being sold to users as defensive mechanisms (Table 2), they also represent identified vulnerabilities by attackers (Table 3). This "defender-as-target" paradox was not created by individual vendors but is an issue with the entire AML industry – there has never been a single vendor-identified model-layer solution for the AML.T0051 threat that has stood up to adversarial testing, and all vendors' products using LLMs have had prompt-injection vulnerabilities disclosed. The various vendor responses to these issues illustrate the conflict they are experiencing: Anthropic, Google, and Microsoft each paid bug bounties for the comment-and-control attack (\$100, \$1,337, and \$500, respectively), but none issued a CVE or a public advisory. Only Anthropic updated its product documentation, warning customers that the affected component is not hardened against prompt injection. For security organizations, the implication is straightforward: any LLM deployed defensively must be documented as a potential attack surface in its own right.

6. Findings and Discussion

6.1. Finding 1: Capability versus Policy Gap

Three findings emerge from the analysis. The first is a capability-versus-policy gap. Across the three frontier LLMs, raw model capability is close enough that benchmark scores alone do not determine real-world utility. As Table 1's cost and accessibility row shows, deployment policy — Anthropic's Cyber Verification Program, OpenAI's Trusted Access for Cyber program, Google's Enterprise tier pricing, and the closed Project Glasswing consortium — determines which model an organization can actually use for which workflow. The implication is that security organizations evaluating LLMs must look past published benchmarks and focus on access-tier requirements. A model that scores HIGH on capability fit but requires a vetting program that the organization cannot enter may, in practice, score effectively zero for that organization.

6.2. Finding 2: Asymmetric Specialization

The second finding is one of asymmetric specialization. Anthropic Claude and OpenAI GPT-5 compete head-to-head on code-centric defensive work; Claude Security and Codex Security map to overlapping use cases in vulnerability discovery, exploit validation, and patch generation. Google Gemini for SecOps does not enter that competition. Its position rests entirely on operational SOC integration: native SIEM and SOAR plumbing, more than 300 prebuilt tool integrations, and a domain-specialized base model trained on security data. Table 1's operational integration row scores Gemini HIGH and the other two MED-HIGH, which reflects this asymmetry. The selection question for an organization is therefore not which LLM is best, but which LLM is best for which workflow, given which budget and which existing tooling.

6.3. Finding 3: The Defender-as-Target Paradox

The third finding is the defender-as-target paradox. Section 5's ATLAS mapping documents prompt-injection attacks against every LLM in the comparison. This is an industry-level weakness rather than a vendor-specific one; no published model-layer defense for AML.T0051 has held up under adversarial testing. While detection-layer responses are maturing through SIEM-side rules mapped to specific ATLAS technique IDs, the underlying technique class remains unsolved at the model layer. Organizations deploying LLMs for defense must therefore plan for the LLM itself to be a monitored asset, not only an analytic tool. Standard security telemetry — input logs, output logs, and tool-invocation traces — should be routed into the SIEM and correlated against ATLAS-mapped detection rules in the same way that endpoint or network telemetry already is.

6.4. Implications for Security Organizations

These findings translate into four recommendations for security organizations. First, match the LLM to the workflow rather than selecting a single vendor across all workflows; Table 2 provides a starting framework. Second, plan for dual deployment in environments that span both code-centric and SOC-centric work — Claude or OpenAI for vulnerability discovery and patching, Google Gemini for SecOps for SIEM-driven detection and response. Third, treat every deployed LLM as a potential attack surface: instrument it with logging, ingest its telemetry into the SIEM, and apply ATLAS-mapped detection rules. Fourth, align governance with the NIST AI RMF Generative AI Profile (NIST AI 600-1), and for DoD-adjacent or critical-infrastructure work, prioritize the data-handling controls strongest in Table 1’s third row — Google’s CMEK and VPC Service Controls, Anthropic’s Cyber Verification Program, and OpenAI’s Trusted Access for Cyber program.

7. Conclusion

7.1. Summary

This paper compared three frontier large language models — Anthropic Claude, OpenAI GPT-5, and Google Gemini for SecOps — across the five evaluation criteria summarized in Table 1 and the MITRE ATT&CK and ATLAS technique mappings presented in Tables 2 and 3. The analysis produced three findings: deployment policy now matters more than benchmark capability for real-world LLM selection; the three vendors exhibit asymmetric specialization, with Anthropic and OpenAI competing on code-centric defensive work while Google competes alone on SOC integration; and every defensively-deployed LLM is itself a documented attack surface subject to industry-wide prompt-injection vulnerabilities. In cybersecurity, frontier LLMs are simultaneously tools and targets, and the role any one of them plays in a given engagement depends as much on the framework being applied as on the model itself.

7.2. Limitations and Future Work

Two limitations are worth addressing. Cyber-focused variants: OpenAI’s GPT-5.4-Cyber and GPT-5.5-Cyber, and Anthropic’s Claude Mythos via Project Glasswing, are restricted to only vetted organizations and have not been independently benchmarked by independent organizations/persons. The empirical claims about these models are therefore vendor attested. Second, the threat landscape documented in Section 2.3 evolves monthly; therefore, incidents catalogued here represent the state of the field through early May 2026 and should not be treated definitively. Future work could extend this analysis in two directions: empirical evaluation of the three LLMs in production SOC environments against a common benchmark, and longitudinal tracking of the field as MITRE ATLAS expands its technique catalog.

References

- Abuadbbba, A., Hicks, C., Moore, K., Mavroudis, V., Hasircioglu, B., Goel, D., & Jennings, P. (2025). From Promise to Peril: Rethinking Cybersecurity Red and Blue Teaming in the Age of LLMs. arXiv preprint arXiv:2506.13434. <https://arxiv.org/abs/2506.13434>
- Anthropic. (2026, April 30). Introducing Claude Security. Retrieved from <https://www.anthropic.com/news>
- Cybersecurity and Infrastructure Security Agency, National Security Agency, & Federal Bureau of Investigation. (2024). Deploying AI Systems Securely. Joint Cybersecurity Information Sheet. <https://www.cisa.gov>
- Google Cloud. (2026). Gemini in Google SecOps Documentation. Retrieved from <https://docs.cloud.google.com/chronicle/docs/secops/gemini-chronicle>
- Guan, A., Fendley, N., Liu, Z., Senapati Diwangkara, & Cao, Y. (2026, April 15). Anthropic, Google, Microsoft paid AI bug bounties — quietly. The Register. https://www.theregister.com/2026/04/15/claude_gemini_copilot_agents_hijacked/
- MITRE Corporation. (2026). MITRE ATT&CK Framework. <https://attack.mitre.org>
- MITRE Corporation. (2026). MITRE ATLAS: Adversarial Threat Landscape for Artificial-Intelligence Systems. <https://atlas.mitre.org>
- National Institute of Standards and Technology. (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0). NIST AI 100-1. <https://www.nist.gov/itl/ai-risk-management-framework>
- National Institute of Standards and Technology. (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile. NIST AI 600-1. <https://airc.nist.gov>
- OpenAI. (2025, October). Introducing Aardvark: OpenAI's Agentic Security Researcher. <https://openai.com/index/introducing-aardvark>
- OpenAI. (2026). Scaling Trusted Access for Cyber Defense. <https://openai.com/index/scaling-trusted-access-for-cyber-defense>
- OWASP Foundation. (2025). OWASP Top 10 for Large Language Model Applications. <https://genai.owasp.org/llm-top-10>

Taler, D. (2026, January 15). Researchers Reveal Reprompt Attack Against Microsoft Copilot. The Hacker News. <https://thehackernews.com/2026/01/researchers-reveal-reprompt-attack.html>